

E-commerce

Sales Performance Analysis using PYTHON and SQL

Introduction

Our company is a rapidly growing e-commerce platform serving customers across Brazil, with a presence in over 4,000 cities and achieving 20% year-over-year growth in 2018. We offer a wide range of products—from Health & Beauty to Electronics—connecting customers with sellers through a user-friendly, installment-friendly shopping experience.

While we've successfully expanded our reach and built a strong customer acquisition engine, our next phase of growth is focused on deepening customer relationships and diversifying our operations. By investing in loyalty programs, geographic expansion beyond São Paulo, and a broader, more inclusive seller network, we aim to transform one-time buyers into repeat customers and create a more resilient, sustainable business model.

With strong market potential and a clear roadmap, we are positioned to become a leading player in Brazil's e-commerce ecosystem.

BASIC PROBLEMS

1. List all unique cities where customers are located.

The screenshot shows a database interface with a toolbar at the top and several tabs labeled "SQL FIELDS", "SQL FIELDS 4", "ADVANCED FIELDS", "INTERMEDIATE FIELDS", and "BASIC FIELDS". Below the toolbar, there is a query editor window containing the following SQL code:

```
1  ## 1. List all unique cities where customers are located.  
2 •  SELECT DISTINCT customer_city FROM Customers;
```

Below the query editor is a "Result Grid" table with a single column labeled "customer_city". The table contains the following data:

customer_city
franca
sao bernardo do campo
sao paulo
mogi das cruzes
campinas
jaragua do sul
timoteo
curitiba
belo horizonte
montes claros
rio de janeiro
lencois paulista
caxias do sul
piracicaba
guardulhos
naraia

At the bottom of the interface, it says "Customers 1 X".

```
[ ] ## 1. List all unique cities where customers are located.  
import pandas as pd  
df_customers = pd.read_csv('Customers.csv')  
unique_cities = df_customers['customer_city'].unique()  
print(unique_cities)
```

→ ['franca' 'sao bernardo do campo' 'sao paulo' ... 'monte bonito'
 'sao rafael' 'eugenio de castro']

2. Count the number of orders placed in 2017.

```
[ ] ## 2. Count the number of orders placed in 2017.  
orders_df = pd.read_csv('orders.csv')  
orders_df['order_purchase_timestamp'] = pd.to_datetime(orders_df['order_purchase_timestamp'])  
orders_2017 = orders_df[orders_df['order_purchase_timestamp'].dt.year == 2017]  
number_of_orders_2017 = len(orders_2017)  
print(f"Number of orders placed in 2017: {number_of_orders_2017}")
```

→ Number of orders placed in 2017: 45101

The screenshot shows a MySQL Workbench interface. At the top, there's a toolbar with various icons for file operations, search, and database management. Below the toolbar is a query editor window containing the following SQL code:

```
7 ## 2.Count the number of orders placed in 2017.  
8 • SELECT count(order_id) as OrderCountin2017 FROM orders WHERE order_purchase_timestamp LIKE '2017%';  
9  
10  
11
```

The number 8 followed by a bullet point indicates the current step or row being executed. The result grid below the editor shows one row of data:

OrderCountin2017
45101

At the bottom of the interface, there are buttons for "Result Grid", "Filter Rows:", "Export:", and "Wrap Cell Content:".

3. Find the total sales per category.

```
## 3. Find the total sales per category.  
order_items_df = pd.read_csv('order_items.csv')  
products_df = pd.read_csv('products.csv')  
merged_df = pd.merge(order_items_df, products_df, on='product_id')  
sales_per_category = merged_df.groupby('product category')['price'].sum().reset_index()  
sales_per_category_sorted = sales_per_category.sort_values(by='price', ascending=False)  
print("Total sales per product category:")  
print(sales_per_category_sorted)
```

Total sales per product category:

	product category	price
30	HEALTH BEAUTY	1258681.34
45	Watches present	1205005.68
49	bed table bath	1036988.68
68	sport leisure	988048.97
53	computer accessories	911954.32
..
58	flowers	1110.04
32	House Comfort 2	760.27
50	cds music dvds	730.00
18	Fashion Children's Clothing	569.85
62	insurance and services	283.29

[73 rows x 2 columns]

```
8  
9     ## 3. Find the total sales per category.  
10 • ALTER TABLE products CHANGE `product category` product_category VARCHAR(255);  
11 • SELECT  
12     t1.product_category,  
13     SUM(t2.price) AS total_sales  
14     FROM products AS t1  
15     JOIN order_items AS t2  
16     ON t1.product_id = t2.product_id  
17     GROUP BY  
18     t1.product_category;  
19
```

Result Grid | Filter Rows: [] | Export: [] | Wrap Cell Content: []

	product_category	total_sales
1	HEALTH BEAUTY	1258681.3399999724
2	sport leisure	988048.9700000428
3	Cool Stuff	635290.8500000002
4	computer accessories	911954.3200000391
5	Watches present	1205005.6799999962
6	housewares	632248.6600000234

Result 3 x

4. Calculate the percentage of orders that were paid in installments.

```
[ ] ## 4. Calculate the percentage of orders that were paid in installments.  
payments_df = pd.read_csv('payments.csv')  
total_orders = payments_df['order_id'].nunique()  
installment_orders_df = payments_df[payments_df['payment_installments'] > 1]  
installment_orders_count = installment_orders_df['order_id'].nunique()  
percentage_with_installments = (installment_orders_count / total_orders) * 100  
print(f"Percentage of orders paid in installments: {percentage_with_installments:.2f}%")
```

→ Percentage of orders paid in installments: 51.46%

```
--  
20      ## 4. Calculate the percentage of orders that were paid in installments.  
21 •   SELECT  
22     ( COUNT(DISTINCT CASE  
23       WHEN payment_installments > 1  
24         THEN T1.order_id  
25       END)  
26     ) * 100.0 / COUNT(DISTINCT T2.order_id) as InstallemtOrder  
27   FROM payments AS T1  
28   JOIN orders AS T2  
29     ON T1.order_id = T2.order_id;  
30  
31
```

InstallemtOrder
51.45817

5. Count the number of customers from each state.

```
[ ] ## 5.Count the number of customers from each state.  
customers_df = pd.read_csv('Customers.csv')  
customers_by_state = customers_df.groupby('customer_state')['customer_id'].count().reset_index()  
customers_by_state.rename(columns={'customer_id': 'number_of_customers'}, inplace=True)  
customers_by_state_sorted = customers_by_state.sort_values(by='number_of_customers', ascending=False)  
print("Number of customers from each state:")  
print(customers_by_state_sorted)
```

→ Number of customers from each state:

	customer_state	number_of_customers
25	SP	41746
18	RJ	12852
10	MG	11635
22	RS	5466
17	PR	5045
23	SC	3637
4	BA	3380
6	DF	2140
7	ES	2033
8	GO	2020
15	PE	1652
5	CE	1336
13	PA	975
12	MT	907
9	MA	747
11	MS	715
14	PB	536
16	PI	495
19	RN	485

```
--  
32     ## 5. Count the number of customers from each state.  
33 •   SELECT  
34     customer_state,  
35     COUNT(customer_id) AS customer_count  
36     FROM Customers  
37     GROUP BY  
38     customer_state  
39     ORDER BY  
40     customer_count DESC;  
41
```

Result Grid | Filter Rows: Export: Wrap Cell Content:

	customer_state	customer_count
▶	SP	41746
	RJ	12852
	MG	11635
	RS	5466
	PR	5045
	SC	3637

Result 5 ×

INTERMEDIATE PROBLEMS

1. Calculate the number of orders per month in 2018.

▶ `## 1.Calculate the number of orders per month in 2018.`

```
orders_df = pd.read_csv('orders.csv')
orders_df['order_purchase_timestamp'] = pd.to_datetime(orders_df['order_purchase_timestamp'])
orders_2018 = orders_df[orders_df['order_purchase_timestamp'].dt.year == 2018]
orders_per_month_2018 = orders_2018.groupby(orders_2018['order_purchase_timestamp'].dt.month)[['order_id']].count().reset_index()
orders_per_month_2018.columns = ['Month', 'Number of Orders']
print("Number of orders per month in 2018:")
print(orders_per_month_2018)
```

▶ Number of orders per month in 2018:

Month	Number of Orders
0	1
1	2
2	3
3	4
4	5
5	6
6	7
7	8
8	9
9	10
0	7269
1	6728
2	7211
3	6939
4	6873
5	6167
6	6292
7	6512
8	16
9	4

1 ## 1. Calculate the number of orders per month in 2018.

2 • SELECT

3 DATE_FORMAT(order_purchase_timestamp, '%Y-%m') AS order_month,

4 COUNT(order_id) AS order_count

5 FROM orders

6 WHERE

7 YEAR(order_purchase_timestamp) = 2018

8 GROUP BY

9 order_month

10 ORDER BY

11 order_month;

12

13

Result Grid | Filter Rows: [] | Export: | Wrap Cell Content:

	order_month	order_count
▶	2018-01	7269
	2018-02	6728
	2018-03	7211
	2018-04	6939
	2018-05	6873
	2018-06	6167

Result 1 ×

2. Find the average number of products per order. grouped by customer city.

Python code:

```
customers_df = pd.read_csv('Customers.csv')

orders_df = pd.read_csv('orders.csv')

order_items_df = pd.read_csv('order_items.csv')

orders_with_customers = pd.merge(orders_df, customers_df, on='customer_id')

merged_df = pd.merge(orders_with_customers, order_items_df, on='order_id')

products_per_order = merged_df.groupby(['customer_city', 'order_id']).size().reset_index(name='number_of_products')

average_products_per_city =
products_per_order.groupby('customer_city')['number_of_products'].mean().reset_index()

average_products_per_city.rename(columns={'number_of_products':
'average_products_per_order'}, inplace=True)

average_products_per_city_sorted =
average_products_per_city.sort_values(by='average_products_per_order',
ascending=False)

print("Average number of products per order, grouped by customer city:")

print(average_products_per_city_sorted)
```

```
→ Average number of products per order, grouped by customer city:
  customer_city  average_products_per_order
2619  padre carvalho          7.0
907   celso ramos           6.5
1154    datas               6.0
756   candido godoi           6.0
2264  matias olimpio          5.0
...
1946     jatauba             1.0
1947     jatoba              1.0
1949      jaua                1.0
1950    jaupaci              1.0
1935     jardim              1.0
[4110 rows x 2 columns]
```

The screenshot shows a database result grid titled "Result Grid". It has two columns: "customer_city" and "avg_products_per_order". The data is sorted by average products per order in descending order. The cities listed are: padre carvalho (7.0000), celso ramos (6.5000), datas (6.0000), candido godoi (6.0000), matias olimpio (5.0000), cidelandia (4.0000), picarra (4.0000), curralinho (4.0000), teixeira soares (4.0000), morro de sao paulo (4.0000), inconfidentes (3.5000), ipua (3.2500), alto paraíso de goias (3.0000), nova esperanca d... (3.0000), capela (3.0000), ouvidor (3.0000), and pedregulho (3.0000). The last row shows a separator line with three dots (...).

	customer_city	avg_products_per_order
▶	padre carvalho	7.0000
	celso ramos	6.5000
	datas	6.0000
	candido godoi	6.0000
	matias olimpio	5.0000
	cidelandia	4.0000
	picarra	4.0000
	curralinho	4.0000
	teixeira soares	4.0000
	morro de sao paulo	4.0000
	inconfidentes	3.5000
	ipua	3.2500
	alto paraíso de goias	3.0000
	nova esperanca d...	3.0000
	capela	3.0000
	ouvidor	3.0000
	pedregulho	3.0000
	...	---

SQL Code :

```
WITH ProductsPerOrder AS ( SELECT order_id, COUNT(product_id) AS product_count
FROM order_items GROUP BY order_id ) SELECT T3.customer_city,
AVG(T1.product_count) AS avg_products_per_order FROM ProductsPerOrder AS T1
JOIN orders AS T2 ON T1.order_id = T2.order_id JOIN Customers AS T3 ON
T2.customer_id = T3.customer_id GROUP BY T3.customer_city ORDER BY
avg_products_per_order DESC;
```

3. Calculate the percentage of total revenue contributed by each product category.

Python code :

```
order_items_df = pd.read_csv('order_items.csv')
products_df = pd.read_csv('products.csv')
merged_df = pd.merge(order_items_df, products_df, on='product_id')
total_revenue = merged_df['price'].sum()
revenue_per_category = merged_df.groupby('product
category')['price'].sum().reset_index()
revenue_per_category['percentage'] = (revenue_per_category['price'] / total_revenue
* 100
revenue_percentage_sorted = revenue_per_category.sort_values(by='percentage',
ascending=False)
print(revenue_percentage_sorted)
```

	product category	price	percentage
30	HEALTH BEAUTY	1258681.34	9.260700
45	Watches present	1205005.68	8.865783
49	bed table bath	1036988.68	7.629605
68	sport leisure	988048.97	7.269533
53	computer accessories	911954.32	6.709669
..
58	flowers	1110.04	0.008167
32	House Comfort 2	760.27	0.005594
50	cds music dvds	730.00	0.005371
18	Fashion Children's Clothing	569.85	0.004193
62	insurance and services	283.29	0.002084

[73 rows x 3 columns]

SQL code:

```
category. WITH CategoryRevenue AS ( SELECT T1.product_category, SUM(T2.price) AS
revenue FROM products AS T1 JOIN order_items AS T2 ON T1.product_id =
T2.product_id GROUP BY T1.product_category ) SELECT product_category, (revenue *
100.0) / ( SELECT SUM(revenue) FROM CategoryRevenue ) AS
percentage_of_total_revenue FROM CategoryRevenue ORDER BY
percentage_of_total_revenue DESC;
```

product_category	percentage_of_total_revenue
HEALTH BEAUTY	9.386017665234476
Watches present	8.985756950355755
bed table bath	7.732850054906726
sport leisure	7.367905435491189
computer accessories	6.800465761578263
Furniture Decoration	5.441856810688871
Cool Stuff	4.7373794710120665
housewares	4.71469378546692
automotive	4.41992841730693
Garden tools	3.6185693399803665
toys	3.6088016817904687
babies	3.0561190120597974
perfumery	2.976283957982953
telephony	2.413596720392118
Furniture office	2.042931668296624
stationary store	1.7221493380098196
PCs	1.662641536320873

4. Identify the correlation between product price and the number of times a product has been purchased.

Python Code:

```
import pandas as pd  
  
order_items_df = pd.read_csv('order_items.csv')  
  
products_df = pd.read_csv('products.csv')  
  
merged_df = pd.merge(order_items_df, products_df, on='product_id', how='inner')  
  
product_summary =  
  
merged_df.groupby('product_id').agg(number_of_purchases=('order_id',  
'count'), average_price=('price', 'mean')).reset_index()  
  
correlation =  
  
product_summary['number_of_purchases'].corr(product_summary['average_price'])  
  
print(f"The correlation between product price and the number of times a product has  
been  
purchased is: {correlation}")
```

```
print(f"The correlation between product price and the number of times a product has been purchased is: {correlation}")
```

→ The correlation between product price and the number of times a product has been purchased is: -0.032139862680945167

SQL Code:

```
WITH ProductsPerOrder AS ( SELECT order_id, COUNT(product_id) AS  
product_count FROM order_items GROUP BY order_id )  
SELECT T3.customer_city,  
AVG(T1.product_count) AS avg_products_per_order  
FROM ProductsPerOrder AS  
T1 JOIN orders AS T2 ON T1.order_id = T2.order_id  
JOIN Customers AS T3 ON  
T2.customer_id = T3.customer_id  
GROUP BY T3.customer_city  
ORDER BY  
avg_products_per_order DESC;
```

Result Grid	
	price_purchase_correlation
▶	-0.03278211803604942

5. Calculate the total revenue generated by each seller and rank them by revenue.

Python Code:

```
order_items_df = pd.read_csv('order_items.csv')

revenue_per_seller =
    order_items_df.groupby('seller_id')['price'].sum().reset_index()
revenue_per_seller.columns = ['seller_id', 'total_revenue']

ranked_sellers = revenue_per_seller.sort_values(by='total_revenue',
                                                ascending=False)

ranked_sellers['rank'] =
    ranked_sellers['total_revenue'].rank(method='dense',
                                         ascending=False).astype(int)

ranked_sellers = ranked_sellers.reset_index(drop=True)

print("\nSellers ranked by total revenue:")

print(ranked_sellers)
```

Sellers ranked by total revenue:

	seller_id	total_revenue	rank
0	4869f7a5dfa277a7dca6462dcf3b52b2	229472.63	1
1	53243585a1d6dc2643021fd1853d8905	222776.05	2
2	4a3ca9315b744ce9f8e9374361493884	200472.92	3
3	fa1c13f2614d7b5c4749cbc52fecda94	194042.03	4
4	7c67e1448b00f6e969d365cea6b010ab	187923.89	5
...
3090	34aefef746cd81b7f3b23253ea28bef39	8.00	2774
3091	702835e4b785b67a084280efca355756	7.60	2775
3092	1fa2d3def6adfa70e58c276bb64fe5bb	6.90	2776
3093	77128dec4bec4878c37ab7d6169d6f26	6.50	2777
3094	cf6f6bc4df3999b9c6440f124fb2f687	3.50	2778

[3095 rows x 3 columns]

SQL Code:

```
SELECT T1.seller_id, SUM(T2.price) AS total_revenue, RANK() OVER (
    ORDER BY SUM(T2.price)
        DESC ) AS revenue_rank FROM sellers AS T1 JOIN order_items AS T2
    ON T1.seller_id =
        T2.seller_id GROUP BY T1.seller_id ORDER BY revenue_rank;
```

The screenshot shows a database result grid with three columns: seller_id, total_revenue, and revenue_rank. The data is sorted by revenue_rank in descending order. The first few rows of data are identical to the Python output above, but the table continues with many more sellers, showing a wide range of total revenues from approximately 3.50 to over 200,000.

	seller_id	total_revenue	revenue_rank
1	4869f7a5dfa277a7dca6462dcf3b52b2	229472.6299999981	1
2	53243585a1d6dc2643021fd1853d8905	222776.0499999952	2
3	4a3ca9315b744ce9f8e9374361493884	200472.9199999949	3
4	fa1c13f2614d7b5c4749cbc52fecda94	194042.02999999846	4
5	7c67e1448b00f6e969d365cea6b010ab	187923.8899999995	5
6	7e93a43ef30c4f03f38b393420bc753a	176431.86999999982	6
7	da8622b14eb17ae2831f4ac5b9dab84a	160236.56999999538	7
8	7a67c85e85bb2ce8582c35f2203ad736	141745.53000000177	8
9	1025f0e2d44d7041d6cf58b6550e0bfa	138968.5499999995	9
10	955fee9216a65b617aa5c0531780ce60	135171.69999999914	10
11	46dc3b2cc0980fb8ec44634e21d2718e	128111.1900000011	11
12	6560211a19b47992c3666cc44a7e94c0	123304.83000000003	12
13	620c87c171fb2a6dd6e8bb4dec959fc6	114774.49999999843	13
14	7d13fcfa15225358621be4086e1eb0964	113628.9699999993	14
15	5dceca129747e92ff8ef7a997dc4f8ca	112155.52999999959	15
16	1f50f920176fa81dab994f9023523100	106939.20999999739	16
17	cc419e0650a3c5ba77189a1882b7556a	104288.42000000263	17

ADVANCED PROBLEMS

1. Calculate the moving average of order values for each customer over their order history.

Python Code:

```
import pandas as pd

orders_df = pd.read_csv('orders.csv')
order_items_df = pd.read_csv('order_items.csv')
customers_df = pd.read_csv('Customers.csv')

order_value_df = pd.merge(order_items_df, orders_df, on='order_id', how='left')

order_value_df = order_value_df.groupby('order_id').agg(order_value=('price', 'sum'), order_purchase_timestamp=('order_purchase_timestamp', 'first')).reset_index()

order_value_df = pd.merge(order_value_df, orders_df[['order_id', 'customer_id']], on='order_id', how='left')

order_value_df = pd.merge(order_value_df, customers_df, on='customer_id', how='left')

order_value_df.dropna(subset=['customer_unique_id'], inplace=True)

order_value_df['order_purchase_timestamp'] = pd.to_datetime(order_value_df['order_purchase_timestamp'])

order_value_df.sort_values(by=['customer_unique_id', 'order_purchase_timestamp'], inplace=True)

order_value_df['moving_average_order_value'] =
order_value_df.groupby('customer_unique_id')['order_value'].transform(lambda x: x.rolling(window=2, min_periods=1).mean())

print("Moving Average of Order Values:")

print(order_value_df[['customer_unique_id', 'order_purchase_timestamp', 'order_value', 'moving_average_order_value']].head(10))
```

Moving Average of Order Values:				
	customer_unique_id	order_purchase_timestamp	order_value	\
87338	0000366f3b9a7992bf8c76cfdf3221e2	2018-05-10 10:56:27	129.90	
20578	0000b849f77a49e4a4ce2b2a4ca5be3f	2018-05-07 11:11:27	18.90	
68939	0000f46a3911fa3c0805444483337064	2017-03-10 21:05:03	69.00	
25028	0000f6ccb0745a6a4b88665a16c9f078	2017-10-12 20:29:41	25.99	
83852	0004aac84e0df4da2b147fcfa70cf8255	2017-11-14 19:45:42	180.00	
23946	0004bd2a26a76fe21f786e4fb0607f	2018-04-05 19:33:16	154.00	
80227	00050ab1314c0e55a6ca13cf7181fecf	2018-04-20 12:57:23	27.99	
26528	00053a61a98854899e70ed204dd4bafe	2018-02-28 11:15:41	382.00	
67031	0005e1862207bf6ccc02e4228effd9a0	2017-03-04 23:32:12	135.00	
641	0005ef4cd20d2893f0d9fdbd94d3c0d97	2018-03-12 15:22:12	104.90	
	moving_average_order_value			
87338		129.90		
20578		18.90		
68939		69.00		
25028		25.99		
83852		180.00		
23946		154.00		
80227		27.99		
26528		382.00		
67031		135.00		
641		104.90		

SQL Code:

```
WITH order_values AS ( SELECT t1.order_id, t1.order_purchase_timestamp,
t2.customer_unique_id, SUM(t3.price) AS total_order_value FROM orders
AS t1 JOIN Customers AS t2 ON t1.customer_id = t2.customer_id JOIN
order_items AS t3 ON t1.order_id = t3.order_id GROUP BY t1.order_id,
t1.order_purchase_timestamp, t2.customer_unique_id)SELECT order_id,
customer_unique_id, order_purchase_timestamp, total_order_value,
AVG(total_order_value) OVER ( PARTITION BY customer_unique_id ORDER BY
order_purchase_timestamp ROWS BETWEEN 1 PRECEDING AND CURRENT ROW )
AS moving_average_order_valueFROM order_valuesORDER BY
customer_unique_id, order_purchase_timestamp;
```

Result Grid					
order_id	customer_unique_id	order_purchase_timestamp	total_order_value	moving_average	Export: Filter Rows: Wrap Cell Content: Fetch rows:
e22acc9c116caa3f2b712bbb380d08e	0000366f3b9a7992bf8c76cfdf3221e2	2018-05-10 10:56:27	129.9	129.9	
3594e005ac04d06a72673270ef9ec9	0000b849f77a49e4a4ce2b2a4ca5be3f	2018-05-07 11:11:27	18.9	18.9	
b33ec3b699337181488304f362a6b734	0000f46a3911fa3c0805444483337064	2017-03-10 21:05:03	69	69	
41272756ecdd9a9ed0180413cc22fb6	0000f6ccb0745a6a4b88665a16c9f078	2017-10-12 20:29:41	25.99	25.99	
d957021f1127559cd947b62533f48f7	0004aac84e0df4da2b147fcfa70cf8255	2017-11-14 19:45:42	180	180	
3e470077b690ea3e3d501cff5e0499	0004bd2a26a76fe21f786e4fb0607f	2018-04-05 19:33:16	154	154	
d0028facea13f50a88020d7097a5a1	00050ab1314c0e55a6ca13cf7181fecf	2018-04-20 12:57:23	27.99	27.99	
44e608f2db00c74a1fe329de4416a4e	00053a61a98854899e70ed204dd4bafe	2018-02-28 11:15:41	382	382	
ae76be74b97bcb0b3e355e60d9a6f9c	0005e1862207bf6ccc02e4228effd9a0	2017-03-04 23:32:12	135	135	
01b330808c5819a6a3cb79b72fb08288	0005ef4cd20d2893f0d9fb9d43c0d97	2018-03-12 15:22:12	104.9	104.9	
6681163e3dab91c549952b2845b20281	0006fdcc98a402fcbe4eb0ee528f6a8d4	2017-07-18 9:23:10	13.9	13.9	
67503374d1fbcbef5a40324f703ff8	00082dbe03e478190aadbea78542e933	2017-11-19 15:22:02	79	79	
85bf8863657bf31006811d45d1cb9	00090324bbad0e934238830371ba0a	2018-03-24 14:44:41	49.95	49.95	
d342d7bad292d68d41d58fb594f431	000949456b182f53c18b68d6babc79c1	2018-04-23 9:55:46	64.89	64.89	

2. Calculate the cumulative sales per month for each year.

Python Code:

```
orders_df = pd.read_csv('orders.csv')

order_items_df = pd.read_csv('order_items.csv')

order_sales = order_items_df.groupby('order_id').agg(total_sales=('price', 'sum')).reset_index()

order_sales_with_date = pd.merge(orders_df[['order_id', 'order_purchase_timestamp']], order_sales, on='order_id')

order_sales_with_date['order_purchase_timestamp'] = pd.to_datetime(order_sales_with_date['order_purchase_timestamp'])

order_sales_with_date['year'] = order_sales_with_date['order_purchase_timestamp'].dt.year

order_sales_with_date['month'] = order_sales_with_date['order_purchase_timestamp'].dt.month

monthly_sales = order_sales_with_date.groupby(['year', 'month']).agg(monthly_sales=('total_sales', 'sum')).reset_index()

monthly_sales = monthly_sales.sort_values(by=['year', 'month'])

monthly_sales['cumulative_sales'] = monthly_sales.groupby('year')['monthly_sales'].cumsum()

print("\nCumulative Sales per Month for Each Year:")
```

Cumulative Sales per Month for Each Year:			
year	month	monthly_sales	cumulative_sales
0 2016	9	267.36	267.36
1 2016	10	49507.66	49775.02
2 2016	12	10.90	49785.92
3 2017	1	120312.87	120312.87
4 2017	2	247303.02	367615.89
5 2017	3	372344.30	740950.19
6 2017	4	359021.20	1101887.39
7 2017	5	506011.14	1607958.56
8 2017	6	433038.60	2040997.16
9 2017	7	498031.48	2539028.64
10 2017	8	573971.68	3113000.32
11 2017	9	624401.69	3737402.01
12 2017	10	664219.43	4401621.44
13 2017	11	1010271.37	5411892.81
14 2017	12	743914.17	6155806.98
15 2018	1	950030.36	950030.36
16 2018	2	844178.71	1794209.07
17 2018	3	983213.44	2777422.51
18 2018	4	996647.75	3774070.26
19 2018	5	996517.68	4770587.94
20 2018	6	865124.31	5635712.25
21 2018	7	895507.22	6531219.47
22 2018	8	854686.33	7385905.80
23 2018	9	145.00	7386050.80

SQL Code:

```
WITH MonthlySales AS ( SELECT  
    DATE_FORMAT(T1.order_purchase_timestamp, '%Y') AS order_year,  
    DATE_FORMAT(T1.order_purchase_timestamp, '%Y-%m') AS order_month,  
    SUM(T2.price) AS monthly_sales FROM orders AS T1 JOIN order_items AS T2  
    ON T1.order_id = T2.order_id GROUP BY order_year, order_month ) SELECT  
    order_year, order_month, monthly_sales, SUM(monthly_sales) OVER (  
    PARTITION BY order_year ORDER BY order_month ) AS cumulative_sales FROM  
    MonthlySales ORDER BY order_year, order_month;
```

order_year	order_month	monthly_sales	cumulative_sales
2016	2016-09	267.36	267.36
2016	2016-10	49507.66000000018	49775.02000000018
2016	2016-12	10.9	49785.92000000018
2017	2017-01	120312.86999999972	120312.86999999972
2017	2017-02	247303.0199999959	367615.8899999956
2017	2017-03	374344.300000005	741960.1899999961
2017	2017-04	359927.2299999999	1101887.419999996
2017	2017-05	506071.1400000094	1607958.5600000054
2017	2017-06	433038.6000000405	2040997.1600000095
2017	2017-07	498031.48000001034	2539028.6400000197
2017	2017-08	573971.6800000152	3113000.3200000348
2017	2017-09	624401.6900000151	3737402.01000005
2017	2017-10	664219.4300000184	4401621.440000068
2017	2017-11	1010271.3700000389	5411892.810000108
2017	2017-12	743914.1700000194	6155806.980000127
2018	2018-01	950030.3600000395	950030.3600000395
2018	2018-02	844178.7100000309	1794209.0700000704

3. Calculate the year-over-year growth rate of total sales.

Python Code:

```
orders_df = pd.read_csv('orders.csv')

order_items_df = pd.read_csv('order_items.csv')

order_sales = order_items_df.groupby('order_id').agg(total_sales=('price', 'sum')).reset_index()

order_sales_with_date = pd.merge(orders_df[['order_id', 'order_purchase_timestamp']], order_sales, on='order_id')

order_sales_with_date['order_purchase_timestamp'] = pd.to_datetime(order_sales_with_date['order_purchase_timestamp'])

order_sales_with_date['year'] = order_sales_with_date['order_purchase_timestamp'].dt.year

annual_sales = order_sales_with_date.groupby('year')['total_sales'].sum().reset_index()

annual_sales['YoY_Growth_Rate'] = annual_sales['total_sales'].pct_change() * 100

print("\nAnnual Sales and Year-over-Year Growth Rate:")

print(annual_sales)

annual_sales.to_csv('year_over_year_growth_rate.csv', index=False)
```

Annual Sales and Year-over-Year Growth Rate:		
year	total_sales	YoY_Growth_Rate
0 2016	49785.92	NaN
1 2017	6155806.98	12264.554035
2 2018	7386050.80	19.985094

SQL Code:

```
SELECT YEAR(T1.order_purchase_timestamp) AS order_year, SUM(T2.price) AS yearly_sales FROM orders AS T1 JOIN order_items AS T2 ON T1.order_id = T2.order_id GROUP BY order_year ) SELECT order_year, yearly_sales, ( yearly_sales - LAG(yearly_sales) OVER ( ORDER BY order_year )) / LAG(yearly_sales) OVER ( ORDER BY order_year ) * 100 AS yoy_growth_rate_percentage FROM AnnualSales ORDER BY order_year;
```

	order_year	yearly_sales	yoy_growth_rate_percentage
▶	2016	49785.92000000019	NULL
	2017	6155806.9800029425	12264.554034560213
	2018	7386050.800004353	19.985094139531032

4. Calculate the retention rate of customers, defined as the percentage of customers who make another purchase within 6 months of their first purchase.

Python code:

```
import pandas as pd

from datetime import timedelta

orders_df = pd.read_csv('orders.csv')

customers_df = pd.read_csv('Customers.csv')

merged_df = pd.merge(orders_df, customers_df, on='customer_id', how='inner')

merged_df['order_purchase_timestamp'] = pd.to_datetime(merged_df['order_purchase_timestamp'])

merged_df.sort_values(by=['customer_unique_id', 'order_purchase_timestamp'], inplace=True)

purchase_dates = merged_df.groupby('customer_unique_id')['order_purchase_timestamp'].apply(list)

retained_customers_data = purchase_dates[purchase_dates.apply(len) > 1].reset_index()

retained_customers_data['first_purchase'] = retained_customers_data['order_purchase_timestamp'].apply(lambda x: x[0])

retained_customers_data['second_purchase'] = retained_customers_data['order_purchase_timestamp'].apply(lambda x: x[1])

retained_customers_data['time_to_second_purchase'] = retained_customers_data['second_purchase'] - retained_customers_data['first_purchase']

six_months = timedelta(days=180)

retained_count = (retained_customers_data['time_to_second_purchase'] <= six_months).sum()

total_retained_customers = len(retained_customers_data)

retention_rate = (retained_count / total_retained_customers) * 100

print(f"Total customers who made at least two purchases: {total_retained_customers}")

print(f"Customers who made a second purchase within 6 months: {retained_count}")

print(f"Customer Retention Rate (within 6 months): {retention_rate:.2f}%")
```

```
Total customers who made at least two purchases: 2997
Customers who made a second purchase within 6 months: 2482
Customer Retention Rate (within 6 months): 82.82%
```

SQL Code:

```
WITH CustomerOrders AS ( SELECT T1.customer_unique_id,
T2.order_purchase_timestamp FROM Customers AS T1 JOIN orders AS T2 ON
T1.customer_id = T2.customer_id ), FirstPurchaseDates AS ( SELECT
customer_unique_id, MIN(order_purchase_timestamp) AS first_purchase_date
FROM CustomerOrders GROUP BY customer_unique_id ), RetainedCustomers AS (
SELECT DISTINCT T1.customer_unique_id FROM FirstPurchaseDates AS T1 JOIN
CustomerOrders AS T2 ON T1.customer_unique_id = T2.customer_unique_id
WHERE T2.order_purchase_timestamp > T1.first_purchase_date AND
T2.order_purchase_timestamp <= DATE_ADD(T1.first_purchase_date, INTERVAL 6
MONTH) ) SELECT ( ( SELECT COUNT(*) FROM RetainedCustomers ) * 100.0 ) / (
SELECT COUNT(*) FROM FirstPurchaseDates ) AS retention_rate_percentage;
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	retention_rate_6_months			
▶	82.98298298298299			

5. Identify the top 3 customers who spent the most money in each year.

Python Code :

```
orders_df = pd.read_csv('orders.csv')

payments_df = pd.read_csv('payments.csv')

df_merged = pd.merge(orders_df, payments_df, on='order_id')

df_merged['order_purchase_timestamp'] = pd.to_datetime(df_merged['order_purchase_timestamp'])

df_merged['year'] = df_merged['order_purchase_timestamp'].dt.year

customer_yearly_spend = df_merged.groupby(['customer_id', 'year'])['payment_value'].sum().reset_index()

customer_yearly_spend.columns = ['customer_id', 'year', 'total_spend']

top_customers_by_year = customer_yearly_spend.groupby('year').apply(lambda x: x.sort_values(by='total_spend', ascending=False).head(3).reset_index(drop=True))

print("Top 3 customers by total spend for each year:")

print(top_customers_by_year)
```

```
print(top_customers_by_year)
```

→ Top 3 customers by total spend for each year:

	customer_id	year	total_spend
0	a9dc96b027d1252bbac0a9b72d837fc6	2016	1423.55
1	1d34ed25963d5aae4cf3d7f3a4cda173	2016	1400.74
2	4a06381959b6670756de02e07b83815f	2016	1227.78
3	1617b1357756262bfa56ab541c47bc16	2017	13664.08
4	c6e2731c5b391845f6800c97401a43a9	2017	6929.31
5	3fd6777bbce08a352fddd04e4a7cc8f6	2017	6726.66
6	ec5b2ba62e574342386871631fafd3fc	2018	7274.88
7	f48d464a0baaea338cb25f816991ab1f	2018	6922.21
8	e0a2412720e9ea4f26c1ac985f6a7358	2018	4809.44

SQL Code :

```
WITH CustomerOrderSpending AS ( SELECT T1.order_id, SUM(T1.payment_value) AS order_value FROM payments AS T1 GROUP BY T1.order_id ), CustomerAnnualSpending AS ( SELECT YEAR(T2.order_purchase_timestamp) AS order_year, T3.customer_unique_id, SUM(T1.order_value) AS total_spent FROM CustomerOrderSpending AS T1 JOIN orders AS T2 ON T1.order_id = T2.order_id JOIN Customers AS T3 ON T2.customer_id = T3.customer_id GROUP BY order_year, T3.customer_unique_id ), RankedCustomers AS ( SELECT order_year, customer_unique_id, total_spent, RANK() OVER ( PARTITION BY order_year ORDER BY total_spent DESC ) AS customer_rank FROM CustomerAnnualSpending ) SELECT order_year, customer_unique_id, total_spent, customer_rank FROM RankedCustomers WHERE customer_rank <= 3 ORDER BY order_year, customer_rank
```

	order_year	customer_unique_id	total_spent	customer_rank
▶	2016	fdaa290acb9eeacb66fa7f979bbaa6803	1423.55	1
	2016	753bc5d6efa9e49a03e34cf521a9e124	1400.74	2
	2016	b92a2e5e8a6eabcc80882c7d68b2c70b	1227.78	3
	2017	0a0a92112bd4c708ca5fd685afaa872	13664.08	1
	2017	da122df9eeaddfedc1dc1f5349a1a690c	7571.63	2
	2017	dc4802a71eae9be1dd28f5d788ceb526	6929.31	3
	2018	46450c74a0d8c5ca9395da1daac6c120	9553.02	1
	2018	763c8b1c9c68a0229c42c9fc6f662b93	7274.88	2
	2018	459bef486812aa25204be022145caa62	6922.21	3

Recommendations

Customer Retention:

Introduce loyalty programs, subscription discounts, or cashback to improve repeat purchase rate.

Personalized product recommendations (using past purchase data) could increase retention.

Geographic Expansion:

Focus marketing on SP, RJ, MG since they contribute the bulk of sales.

But also design growth campaigns in RS, PR, SC where customer base is decent, but sales are relatively lower → untapped potential.

Product Strategy:

Push top categories (Health & Beauty, Watches, Bed/Bath) with targeted ads.

Bundle high-margin but less frequent categories (like electronics, computer accessories) with popular categories.

Payment Optimization:

Since ~50% of payments are installment-based, offering more flexible EMI options could drive higher order values.

Partnerships with banks/fintech's can reduce payment friction.

Seller Network:

Dependency is high on a few sellers. Encourage new seller onboarding from underrepresented states (e.g., Northern Brazil) to diversify.

Provide seller analytics dashboards to improve their sales performance.

Conclusion

The business shows strong YoY growth (20% in 2018) and a wide customer reach (4,119 cities) but suffers from very low customer retention (0%). Sales are highly concentrated in SP state and a few top sellers, which creates risk. To ensure long-term sustainability, the company should:

Build loyalty programs

Expand beyond SP into secondary states

Diversify seller & product portfolio.

With these steps, the platform can convert its large new customer base into repeat loyal buyers, ensuring sustainable revenue growth.